

IPv6 Protocol Architecture

New Functional Improvement

- Address Space
 - Increase from 32-bit to 128-bit address space
- Management
 - StateLessAddress AutoConfiguration(SLAAC) means no more need to configure IP addresses for end systems, even via DHCP
- Performance
 - Simplified header means efficient packet processing
 - No header checksum re-calculation at every hop (when TTL is decremented) => **left up to the lower and upper layers!**
- No hop-by-hop fragmentation -PMTUD

IPv4/IPv6 Header Comparison

Version	IHL	Type of Service	Total Length		Version	Traffic Class	Flow Label						
Identification			Flags	Fragment offset	Payload Length		Next Header	Hop Limit					
Time to Live	Protocol		Header Checksum			Source Address							
Source Address													
Destination Address													
Options				Padding									

Not kept in IPv6

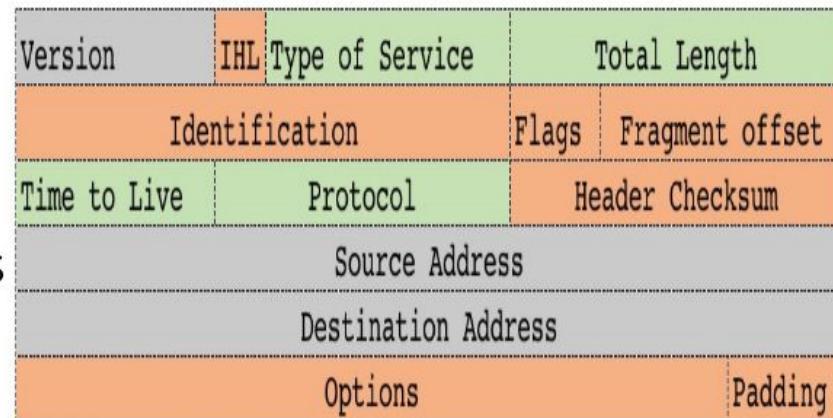
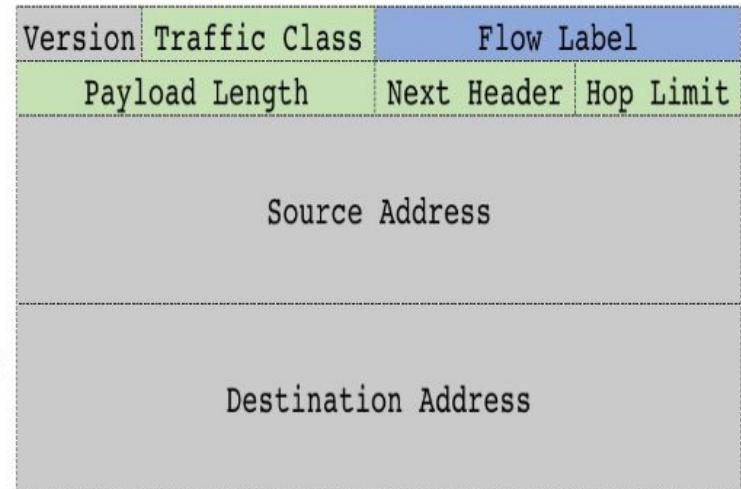
Renamed in IPv6

Same name and function

New in IPv6

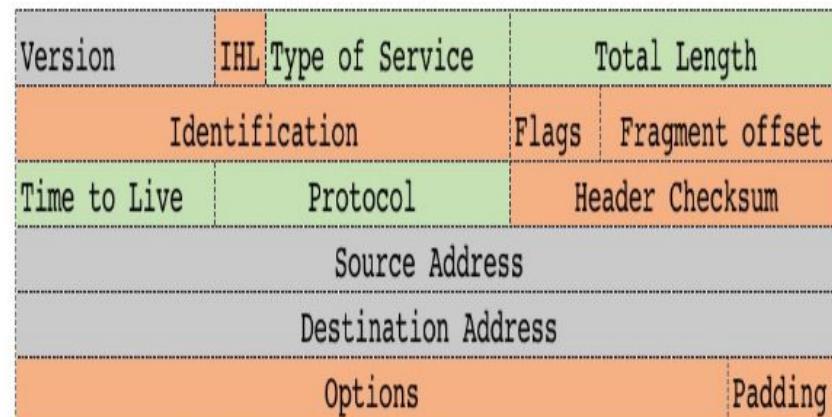
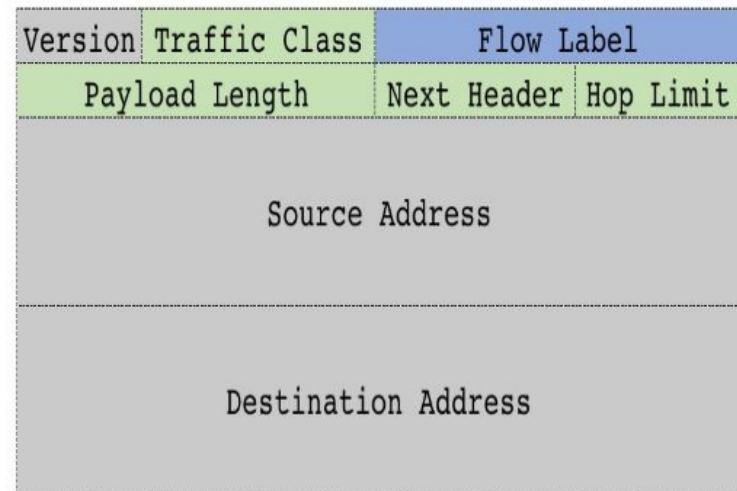
IPv6 Protocol Header Format

- Version (4-bit):
 - 4-bit IP version number (6)
- Traffic class (8-bit):
 - Similar to DiffServ in IPv4; define different classes or priorities.
- Flow label (20-bit):
 - allows IPv6 packets to be identified based on flows (multilayer switching techniques and faster packet-switching performance)



IPv6 Protocol Header Format

- Payload length (16-bit):
 - Defines the length of the IPv6 payload (including extension headers); Total Length in IPv4 includes the header.
- Next header (8-bit):
 - Identifies the type of information following IPv6 header. Could be upper layer (TCP/UDP), or an extension header (similar to Protocol field in IPv4).
- Hop limit (8-bit):
 - Similar to TTL in IPv4

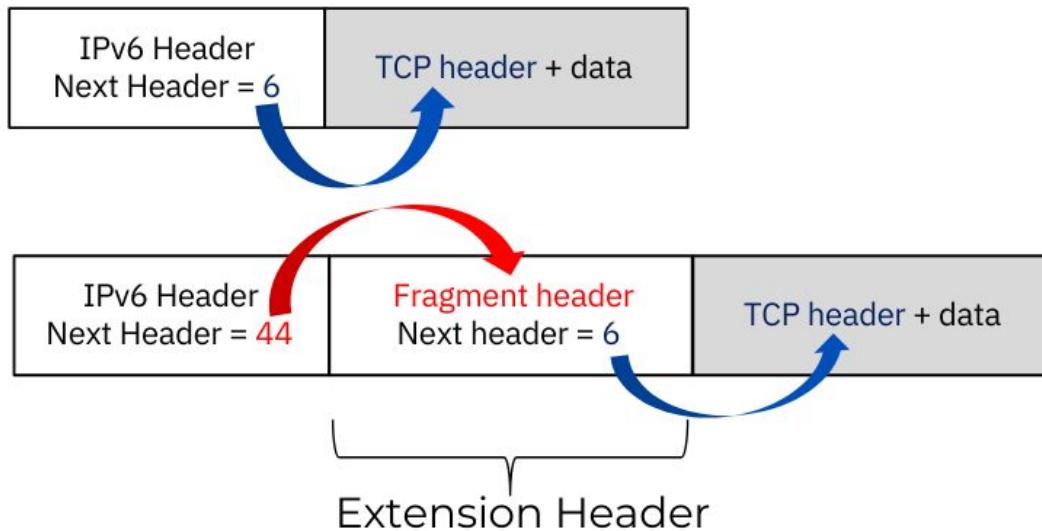


IPv6 & IPv4 Packet Example

- Example IPv6 packet on this link:
<https://www.cloudshark.org/captures/84fd54ad03e0>
- Example IPv4 packet on this link:
<https://www.cloudshark.org/captures/09f49cda5b80>

IPv6 Extension Header

- IPv6 allows an optional *Extension Header* in between the IPv6 header and upper layer header
 - Allows adding new features to IPv6 protocol without major re-engineering



Next Header values:

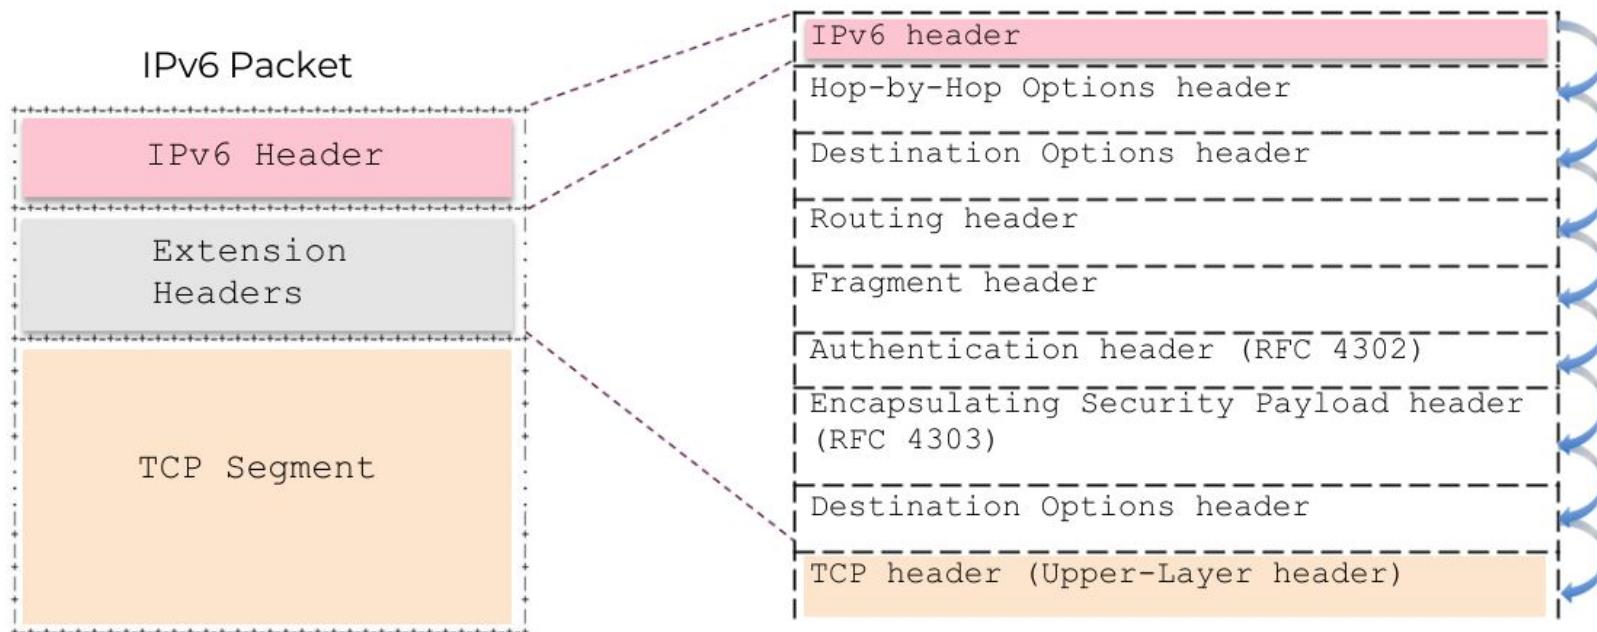
0	Hop-by-hop option
6	TCP
17	UDP
43	Source routing (RFC5095)
44	Fragmentation
50	Encrypted security payload
51	Authentication
58	ICMPv6
59	Null (No next header)
60	Destination option

IPv6 Extension Header (contd)

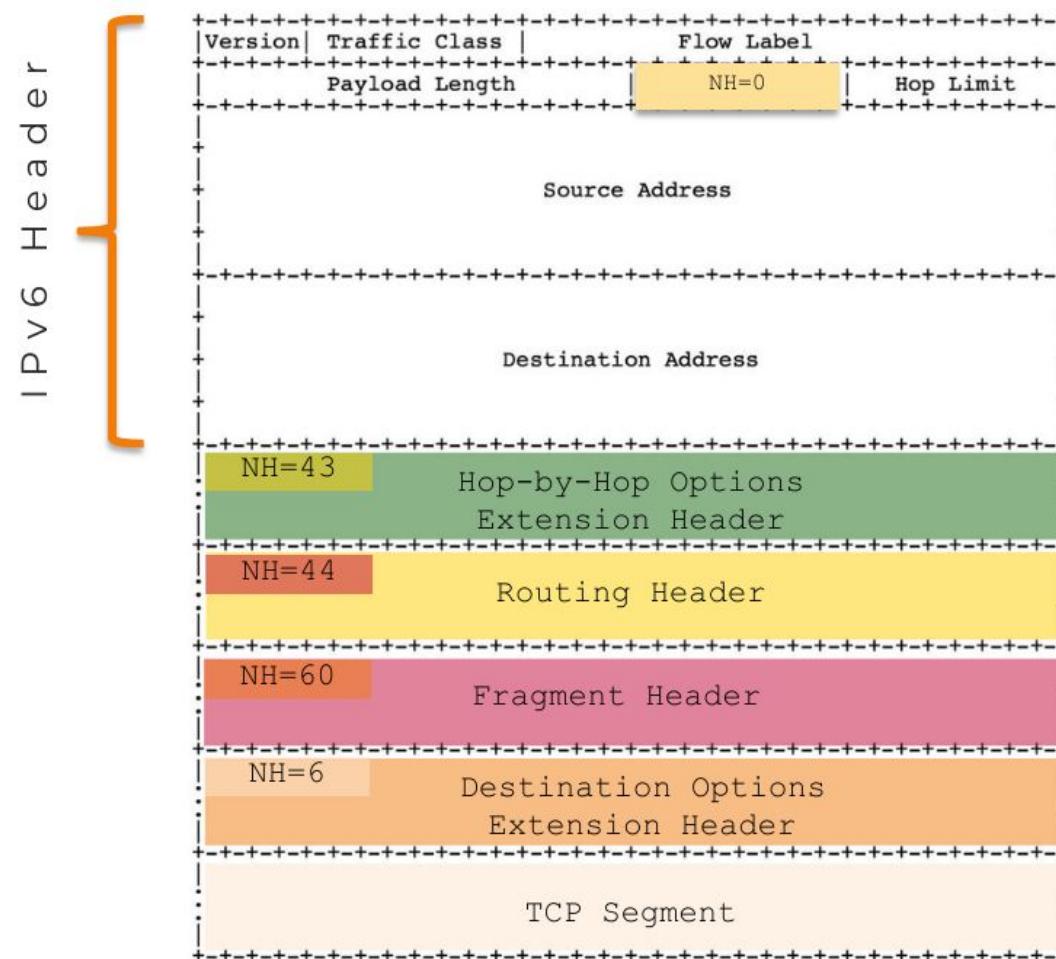
- An IPv6 packet may carry none or many extension headers
 - A next header value of 6 or 17 (TCP/UDP) indicates there is no extension header
 - the next header field points to TCP/UDP header, which is the payload
- Unless the next header value is 0 (*Hop-by-Hop option*), extension headers are processed only by the destination node, specified by the destination address.

Extension Header Order

- When more than one extension header is used in the same packet, it is **recommended** that those headers appear in the following order in RFC 8200:



Chaining Extension Headers



Extension Header Example

- Example IPv6 packet with an Extension Header on this link:
 - <https://www.cloudshark.org/captures/7dd0b50eb768>

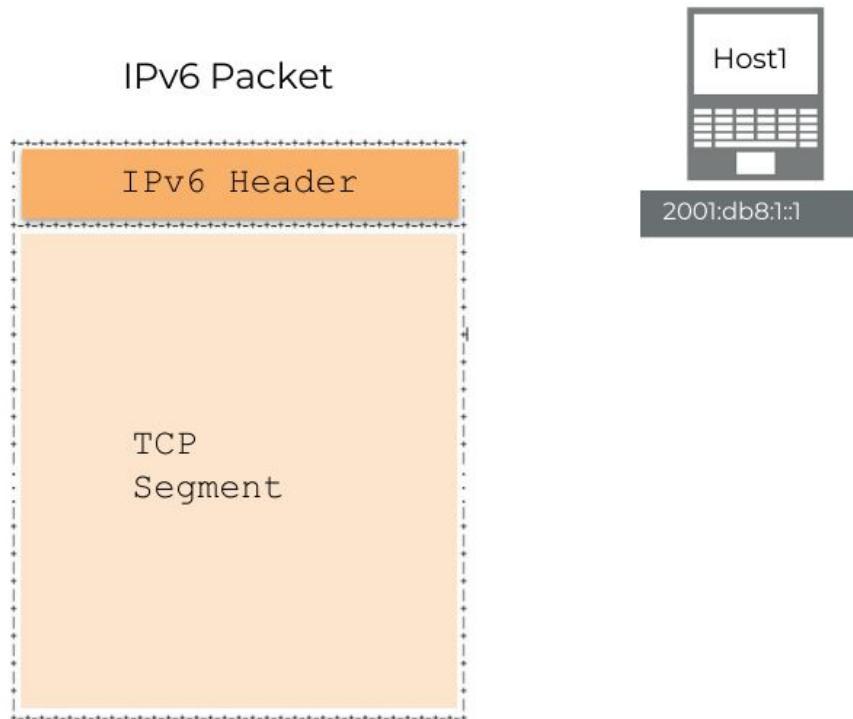
Fragmentation Handling In IPv6

- In IPv6, fragmentation is only performed by the host/source nodes, and not the routers along the path (unlike IPv4)
- Each source device tracks the MTU size for each session
- When a IPv6 host has large amount of data to be sent, it will be send in a series of IPv6 packets (fragmented)
 - IPv6 hosts use Path MTU Discovery (PMTUD) to determine the most optimum MTU size along the path

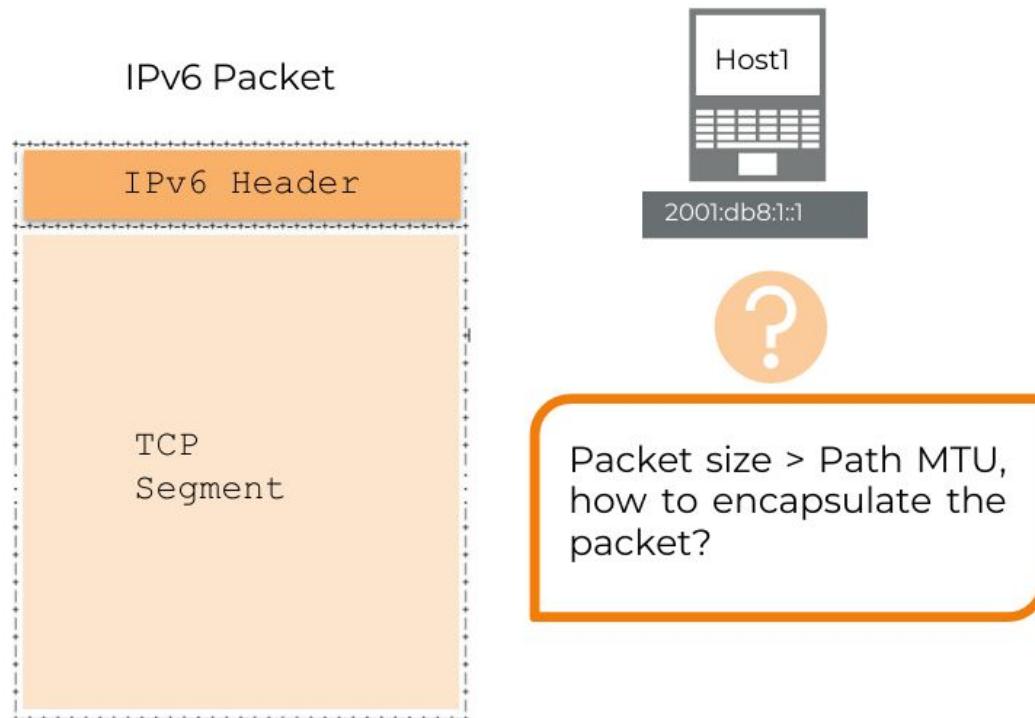
Example of Fragment Header



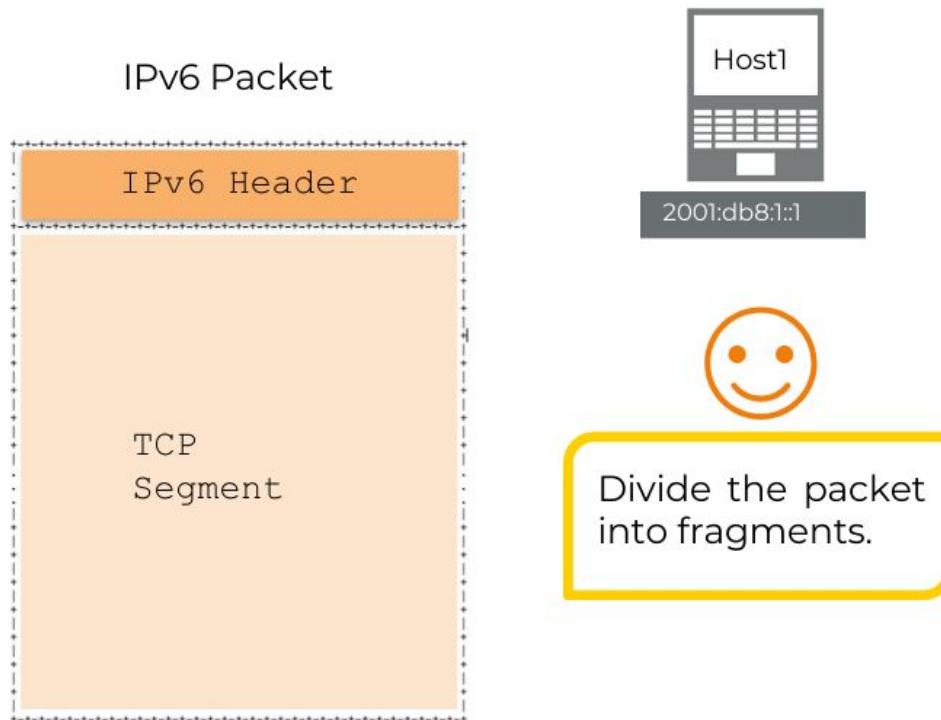
On the Source Node (Host1)



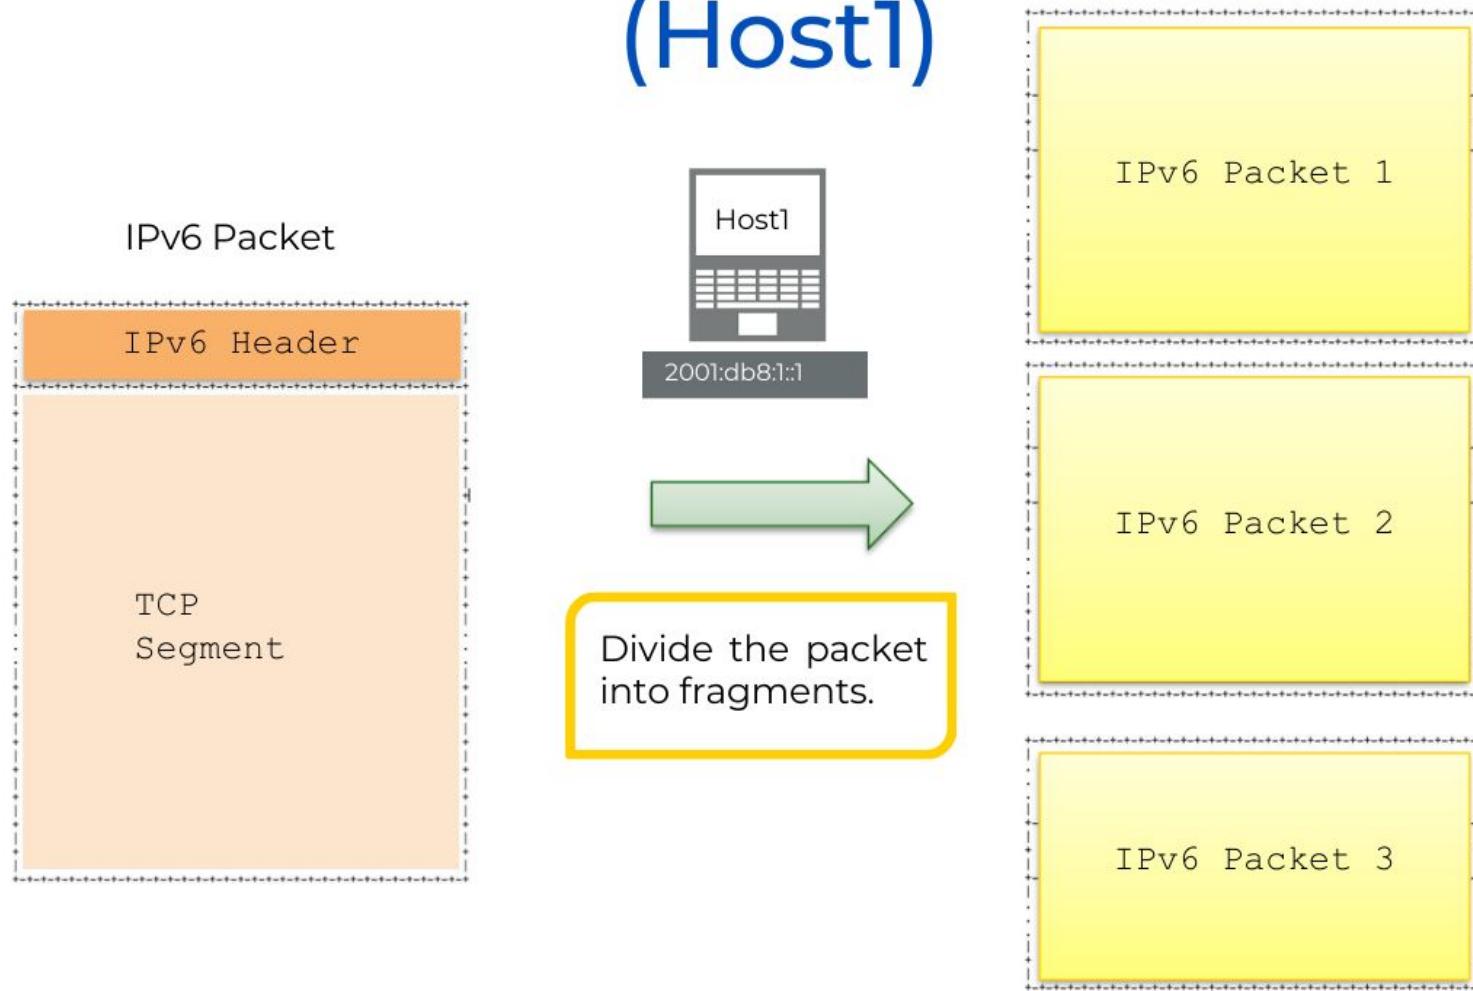
On the Source Node (Host1)



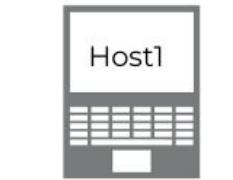
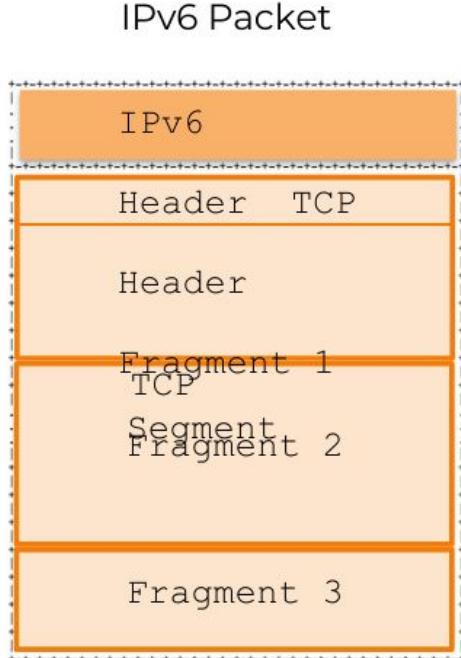
On the Source Node (Host1)



On the Source Node (Host1)



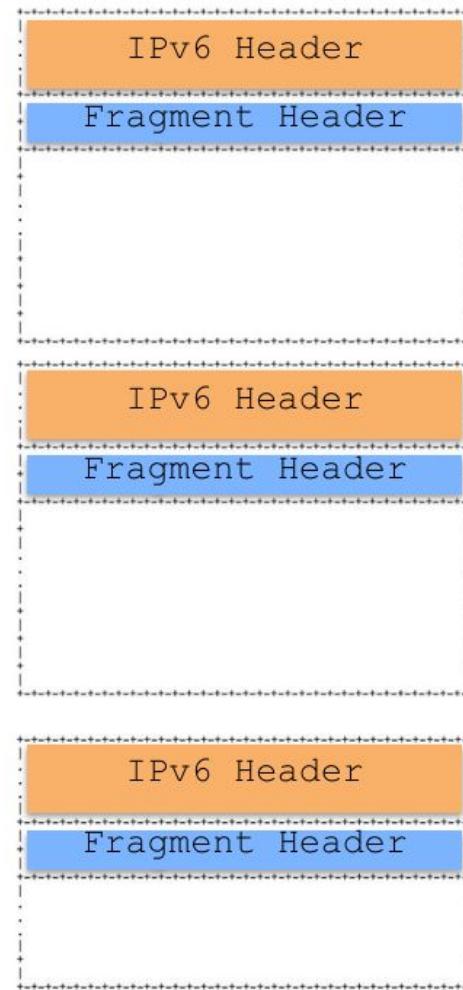
On the Source Node (Host1)



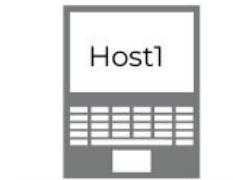
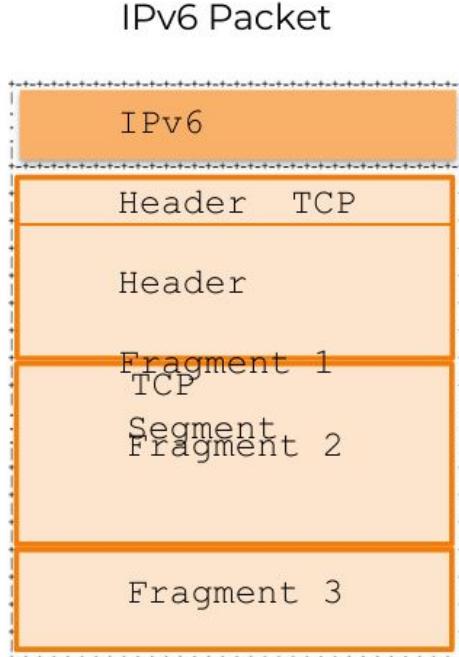
2001:db8:1::1



Divide the packet
into fragments.



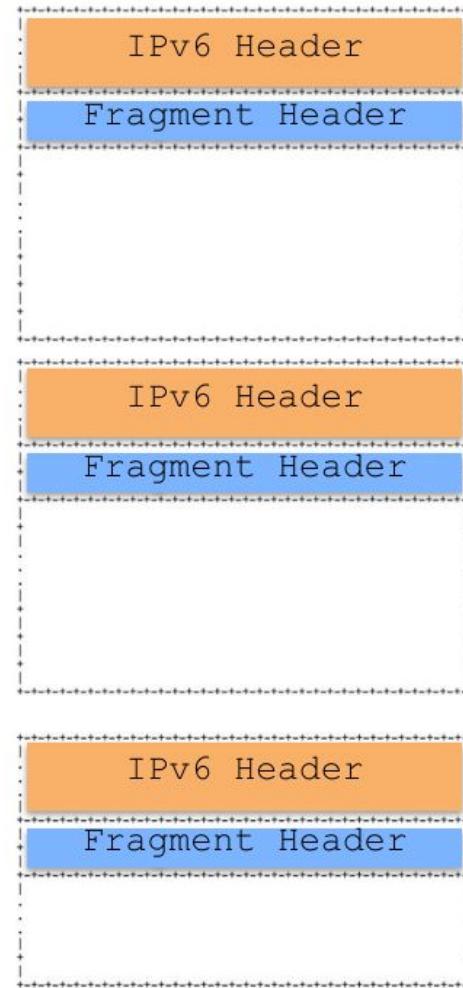
On the Source Node (Host1)



2001:db8:1::1

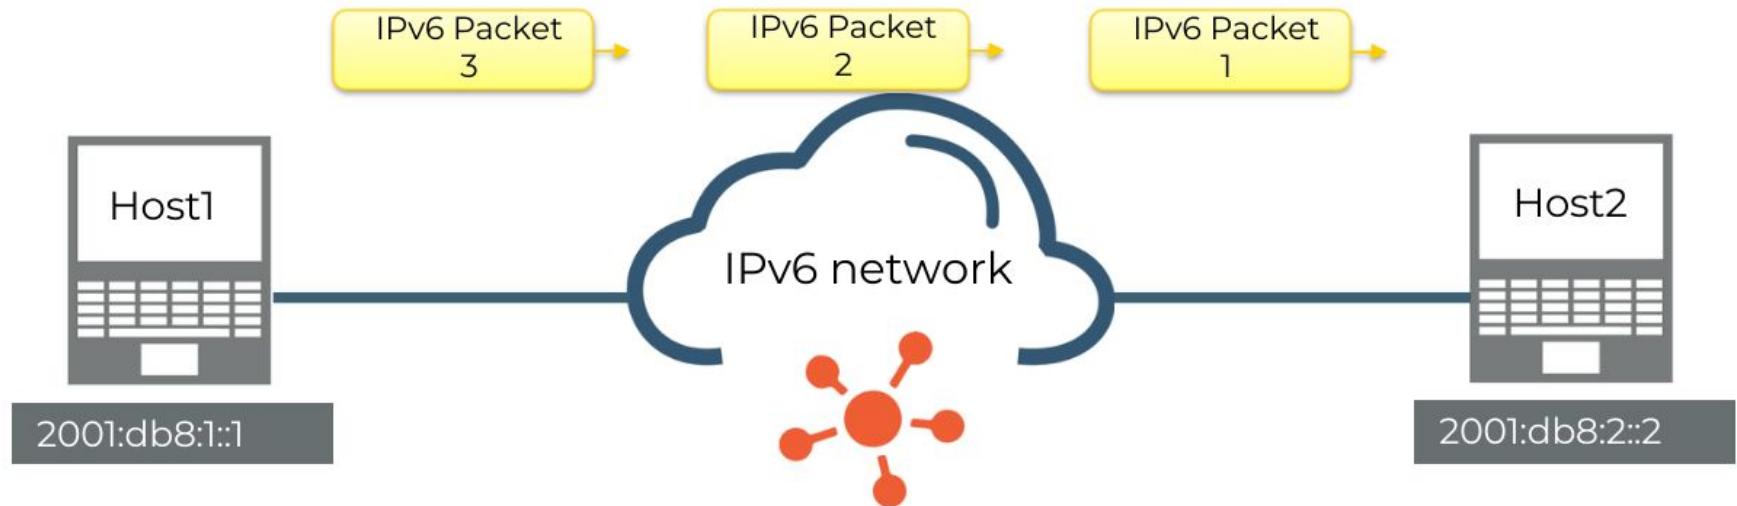


Divide the packet
into fragments.



IPv6 Packet 1
IPv6 Packet 2
IPv6 Packet 3

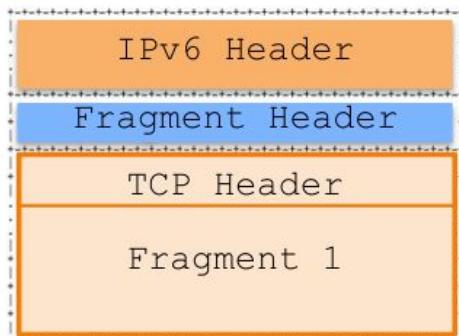
Example of Fragment Header



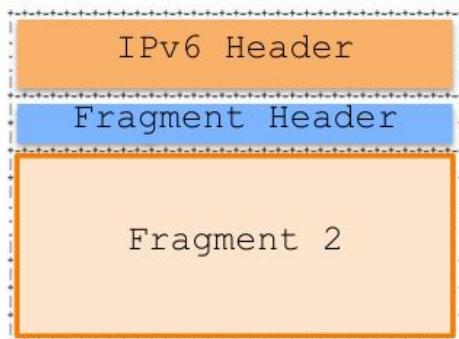
The 3 fragmented packets are transmitted on the path, reach the destination Host2, without any other fragmentation on the path.

On the Destination Node (Host2)

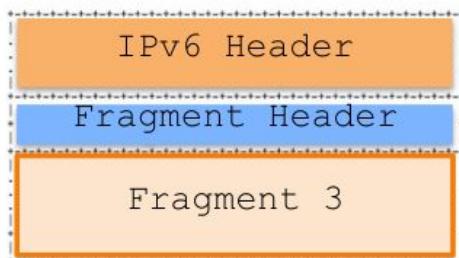
IPv6 Packet 1



IPv6 Packet 2

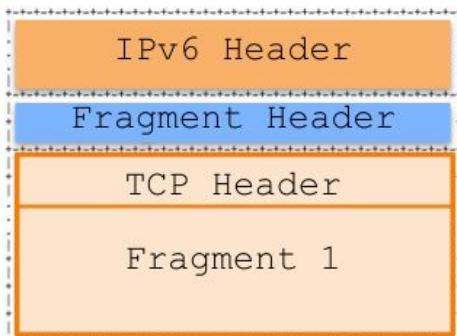


IPv6 Packet 3

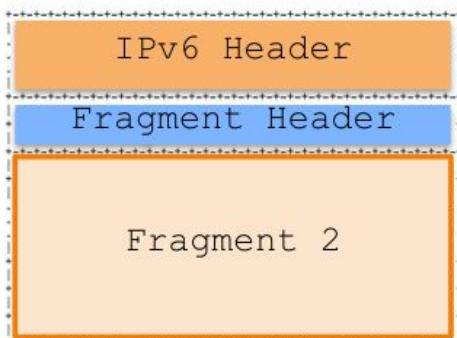


On the Destination Node (Host2)

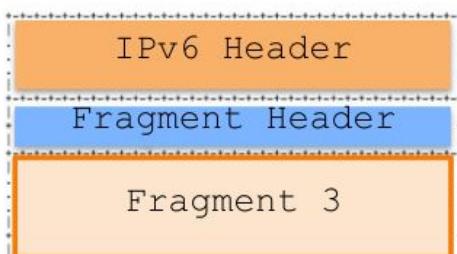
IPv6 Packet 1



IPv6 Packet 2



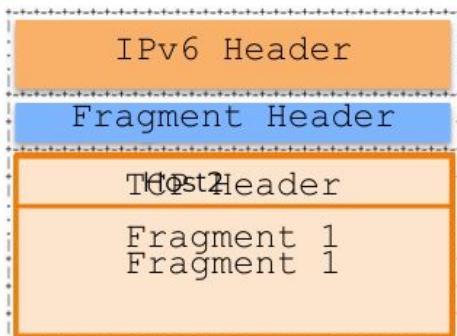
IPv6 Packet 3



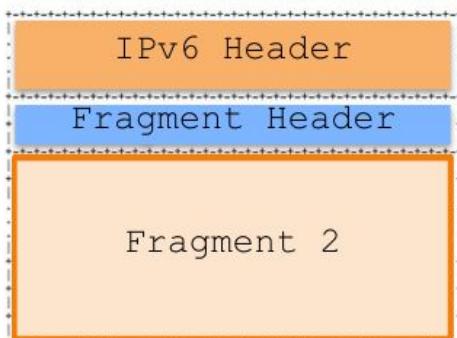
Reassemble the
fragments to be the
original packet.

On the Destination Node (Host2)

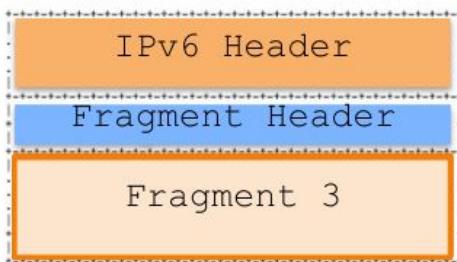
IPv6 Packet 1



IPv6 Packet 2



IPv6 Packet 3

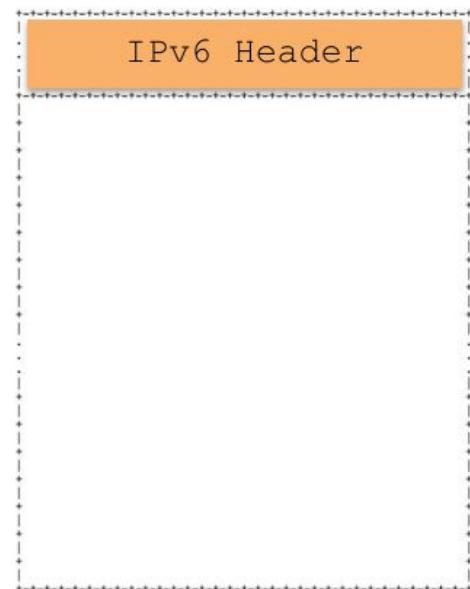


2001:db8:2:2



Reassemble the
fragments to be the
original packet.

IPv6 Packet



On the Destination Node (Host2)

Host2



2001:db8:2:2



Reassemble the
fragments to be the
original packet.

IPv6 Packet

IPv6 Header

TCP
Segment

Path MTU Discovery

- With PMTUD, the source IPv6 device assumes the initial PMTU is the MTU of the first hop in the path
 - upper layers (Transport/Application) send packets based on the first hop MTU
 - If the device receives an “*ICMPv6 packet too big (Type 2)*” message, it informs the upper layer to reduce its packet size, based on the actual MTU size (contained in the message) of the node that dropped the packet

Path MTU Discovery



Link MTU values are marked on each link.

I have a packet with size 2000 bytes to send to Host2. It is larger than MTU, I have to fragment it.

Host1 :

MTU
cache=1500

Path MTU Discovery



Host1 :
MTU cache= **1500**

Link MTU values are marked on each link.

Packet 1
size=1500bytes

Path MTU Discovery



Host1:
MTU cache= **1500**

Link MTU values are marked on each link.

Packet 1
size=1500bytes

```
Frame 1: 1510 bytes on wire (12080 bits), 1510 bytes captured (12080 bits)
Ethernet II, Src: Vmware_23:16:87 (00:0c:29:23:16:87), Dst: Vmware_25:cf:a1 (00:0c:29:25:cf:a1)
Internet Protocol Version 6, Src: 2001:db8:1::1, Dst: 2001:db8:2::2
    0110 .... = Version: 6
    .... 0000 0000 .... .... .... .... = Traffic Class: 0x00 (DSCP: CS0, ECN: Not-ECT)
    .... .... 1010 1000 1001 1111 1000 = Flow Label: 0xa89f8
    Payload Length: 1456
    Next Header: Fragment Header for IPv6 (44)
    Hop Limit: 64
    Source: 2001:db8:1::1
    Destination: 2001:db8:2::2
        Fragment Header for IPv6
    [data]
```

Captured packets are available:
<https://www.cloudshark.org/captures/7dd0b50eb768>

Path MTU Discovery



Link MTU values are marked on each link.

Packet 1
size=1500bytes

Path MTU Discovery



Link MTU values are marked on each link.

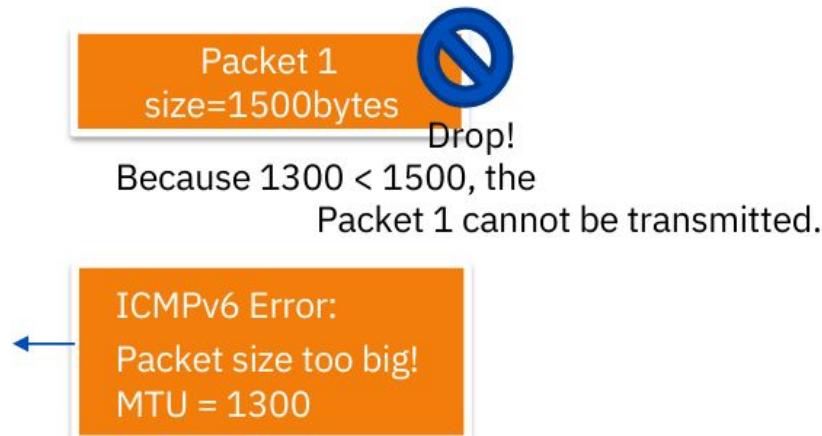
Packet 1
size=1500bytes

Because $1300 < 1500$, the packet 1 cannot be transmitted.

Path MTU Discovery



Link MTU values are marked on each link.



Path MTU Discovery



Link MTU values are marked on each link.

Packet 1
size=1500bytes
Drop!



Because $1300 < 1500$, the packet 1 cannot be transmitted.

ICMP Error:
Packet size too big!
MTU = 1300



Host1 Update :
MTU cache= **1300**

Path MTU Discovery



```
▶ Frame 3: 1294 bytes on wire (10352 bits), 1294 bytes captured (10352 bits)
  Link MTU values are marked on each link.
  ▶ Ethernet II, Src: Vmware_25:cf:a1 (00:0c:29:25:cf:a1), Dst: Vmware_23:16:87 (00:0c:29:23:16:87)
  ▶ Internet Protocol Version 6, Src: 2001:db8:12::2, Dst: 2001:db8:1::1
  ▶ Internet Control Message Protocol v6
    Type: Packet Too Big (2)
    Code: 0
    Checksum: 0x2e57 [correct]
    Checksum Status: Good
    MTU: 1300
    Drroopp!!
    ▶ Internet Protocol Version 6, Src: 2001:db8:1::1, Dst: 2001:db8:2::2
    ▶ data
    Beeccaauussee 11330000 << 11550000,, tthhee ppaacckkeett
    1
    1 ccaannnnoott bbee ttrraannssmmiittteedd..
  ICMPv6 Error:
  Packet size too big!
  MTU = 1300
```

Captured packets are available:
<https://www.cloudshark.org/captures/7dd0b50eb768>

Path MTU Discovery



Host1 :
MTU cache= **1300**

Link MTU values are marked on each link.

Packet 2
size=1300bytes

Path MTU Discovery



Host1:
MTU cache= **1300**

Link MTU values are marked on each link.

Packet 2
size=1300bytes

```
▷ Frame 4: 1310 bytes on wire (10480 bits), 1310 bytes captured (10480 bits)
▷ Ethernet II, Src: Vmware_23:16:87 (00:0c:29:23:16:87), Dst: Vmware_25:cf:a1 (00:0c:29:25:cf:a1)
▽ Internet Protocol Version 6, Src: 2001:db8:1::1, Dst: 2001:db8:2::2
    0110 .... = Version: 6
    .... 0000 0000 .... .... .... .... = Traffic Class: 0x00 (DSCP: CS0, ECN: Not-ECT)
    .... .... .... 1010 1000 1001 1111 1000 = Flow Label: 0xa89f8
    Payload Length: 1256
    Next Header: Fragment Header for IPv6 (44)
    Hop Limit: 64
    Source: 2001:db8:1::1
    Destination: 2001:db8:2::2
    ▷ Fragment Header for IPv6
    ▷ [data]
```

Path MTU Discovery



Host1:
MTU cache=1300

Link MTU values are marked on each link.

Packet 2
size=1300bytes

Path MTU Discovery



Host1 :
MTU cache=1300

Link MTU values are marked on each link.

Packet 2
size=1300bytes

Path MTU Discovery



Host1:
MTU cache=1300

Link MTU values are marked on each link.

Packet 2
size=1300bytes

Path MTU Discovery



Host1:
MTU cache=1300

Link MTU values are marked on each link.

Packet 2
size=1300bytes

Path MTU = 1300



Questions

