



Multihoming Techniques

Customer's Expectation



But it's really just...



Until this happens




Why Multihome?

- Redundancy
 - Avoid a circuit down taking your business down
 - Multiple ways of doing this
- Performance
 - Maybe provider A's connectivity is different to provider B's
 - Asia focus, Europe focus
 - Satellite versus Fiber
- Money
 - Good, Fast, Cheap – Pick Two
 - Buy transit but also peer where you can


Achieving Redundancy

- More than one path to the same ISP
 - Dual-homed




Achieving Redundancy - Multihoming

- More than one upstream ISP
 - Multi-homed




Multihoming – with peering

- One upstream and local peering




Multihoming

- More than one upstream ISP and local peering



Multihoming

- More than one upstream ISP with local and public peering



Traffic Engineering – Path Control

- Remember
 - the prefixes YOU ACCEPT control how YOU SEND traffic
 - Prefixes YOU ANNOUNCE can influence the (path) INCOMING traffic to you
 - There is no guarantee as the other party has control over sending traffic towards you

Traffic Engineering – Accept

- Just because someone sends you a prefix doesn't mean you have to accept it
- Use ***Local Preference*** to control the best path of your outgoing traffic
 - it is the first BGP selection criterion!

Define Local Preference rules

- Follow the money
 - Prefer to send traffic to a customer (earns money)
 - If the route isn't from a customer prefer an announcement seen from a peer (free)
 - Else send to a transit (costs money)
 - Might have a preference between transit providers
- Local preference values should follow this plan
 - Customer > Peer > Transit

Traffic Engineering – Announce


- You can use AS path prepending to try to influence your neighbour's best path
 - but they may be using Local Preference too
- Often, more specific aggregates can help
 - announce them to the preferred path but not others
 - Remember to send the aggregate to everyone
 - Tag the more specific aggregates with the "no-export" community if possible
 - Be careful that you only do this to YOUR prefixes!
- Use MED when there are multiple links to the same AS
 - to signal a preferred path into your AS (may be ignored without prior negotiation)

TE (Path control) Attributes

- Inbound Traffic:
 - AS-PATH, MED, Community (sub-aggregates/more specifics)
- Outbound Traffic:
 - Local Preference

Two Upstream – One backup

- Both incoming and outgoing traffic via AS20 (R1)
- AS30 (R2) path to be used only if the link to AS20 fails
 - AS_PATH to control inbound traffic
 - *Prepend outbound on R2*
 - LOCAL_PREF for outgoing traffic
 - *Higher for inbound routes on R1*



Two Upstream – One backup

- **Always** announce the aggregate on both links!
- **R1** (main link) config:

```
router bgp 17821
 network 61.45.248.0 mask 255.255.248.0
 neighbor 20.20.20.1 remote-as 20
 neighbor 20.20.20.1 prefix-list AGGR out
 neighbor 20.20.20.1 prefix-list DEF in
!
ip prefix-list AGGR permit 61.45.248.0/21
ip prefix-list DEF permit 0.0.0.0/0
!
ip route 61.45.248.0 255.255.248.0 null0
```

Advertise aggregate

Prefix-list applied to
outbound routes

Prefix-list applied to
inbound routes

Define the prefix-lists

Aggregate should exist
in the routing table
(pull-up route)

Two Upstream – One backup

- **R2** (backup) config:

```
router bgp 17821
 network 61.45.248.0 mask 255.255.248.0
 neighbor 30.30.30.1 remote-as 30
 neighbor 30.30.30.1 prefix-list AGGR out
 neighbor 30.30.30.1 route-map BACKUP-OUT out
 neighbor 30.30.30.1 prefix-list DEF in
 neighbor 30.30.30.1 route-map BACKUP-IN in
!
ip prefix-list AGGR permit 61.45.248.0/21
ip prefix-list DEF permit 0.0.0.0/0
!
ip route 61.45.248.0 255.255.248.0 null0
!
route-map BACKUP-OUT permit 10
 set as-path prepend 17821 17821 17821
!
route-map BACKUP-IN permit 10
 set local-preference 80
```

Advertise aggregate in BGP

Route-map applied to
outbound routes

Route-map applied to
inbound routes


Define the prefix-lists

BACKUP-OUT prepends the
AS-PATH for all outbound
BGP updates

BACKUP-IN sets lower local
pref for all inbound BGP
updates

Two Upstream – Load Sharing (Inbound Traffic)

- Announce one sub-aggregate on first, and the other on the second link
 - **Always** announce the aggregate to both!
- Requires good address planning
 - Customers need to be assigned from both address blocks



Two Upstream – Load Sharing (Inbound Traffic)

- R1 config:

```
router bgp 17821
network 61.45.248.0 mask 255.255.248.0
network 61.45.248.0 mask 255.255.252.0
neighbor 20.20.20.1 remote-as 20
neighbor 20.20.20.1 prefix-list SUB-A out

!
ip prefix-list SUB-A permit 61.45.248.0/21
ip prefix-list SUB-A permit 61.45.248.0/22

!
ip route 61.45.248.0 255.255.248.0 null0
ip route 61.45.248.0 255.255.252.0 null0
```

Advertise both the aggregate and first sub-aggregate in BGP

Advertise sub-aggregate along with the aggregate

Sub-aggregate should exist in the routing table (pull-up route)

Two Upstream – Load Sharing (Inbound Traffic)

- R2 config:

```
router bgp 17821
network 61.45.248.0 mask 255.255.248.0
network 61.45.252.0 mask 255.255.252.0
neighbor 30.30.30.1 remote-as 30
neighbor 30.30.30.1 prefix-list SUB-B out

!
ip prefix-list SUB-B permit 61.45.248.0/21
ip prefix-list SUB-B permit 61.45.252.0/22

!
ip route 61.45.248.0 255.255.248.0 null0
ip route 61.45.252.0 255.255.252.0 null0
```

Advertise both
aggregate and second
sub-prefix in BGP


Advertise sub-
aggregate along with
the aggregate

Sub-aggregate should
exist in the routing
table (pull-up route)

Load Sharing – Outbound traffic (Full routes)

- Full Internet routes (**more memory/CPU**)
- Accept full route from both (AS20 and AS30)
 - For routes from AS30 (R2)
 - **Higher LOCAL_PREF** prefixes originated by AS30 and its immediate neighbors (one AS hop away) – traffic to those go via AS30
 - **Lower LOCAL_PREF** all other routes – traffic to these go via AS20
 - For routes learned from AS20 (R1)
 - default LOCAL_PREF value

Load Sharing – Outbound traffic (Full routes)



Load Sharing – Outbound traffic (Full routes)

- R1 configuration: nothing doing

```
router bgp 17821
  neighbor 20.20.20.1 remote-as 20
  neighbor 20.20.20.1 prefix-list ALL in
  !
ip prefix-list ALL deny <bogons>
ip prefix-list ALL permit 0.0.0.0/0 le 24
!
```

Accept full internet feed
except bogons
(default LOCAL_PREF)

Load Sharing – Outbound traffic (Full routes)

- R2 config:

```
router bgp 17821
neighbor 30.30.30.1 remote-as 30
!
address-family ipv4
  neighbor 30.30.30.1 prefix-list ALL in
  neighbor 30.30.30.1 route-map TWO-HOPS in
!
ip prefix-list ALL deny <bogons>
ip prefix-list ALL permit 0.0.0.0/0 le 24
!
ip as-path access-list 30 permit ^(30_)+$
ip as-path access-list 30 permit ^(30_)+_[0-9]+$
!
route-map TWO-HOPS permit 10
  match as-path 30
  set local-preference 150
route-map TWO-HOPS permit 20
  set local-preference 50
```

Accept full internet feed
except bogons

Accept routes local to and
received from AS30
(AS-path prepend included)

Only routes from AS30 and its
direct neighbor ASes


High-pref AS30 and its neighbor
AS originated routes

Low-pref everything else

Load Sharing – Outbound Traffic (Partial routes)

- Partial Routes – **less HW resources!**
- Select some routes for special treatment
- Accept only *default* from AS20
- Default plus routes from AS30 (**better connected than AS20**)
 - Filter to only accept prefixes originated by AS30 and its neighbor ASes (**AS-Path ACLs**)
 - Higher LOCAL_PREF those routes
 - Low LOCAL_PREF the default route
 - Traffic to “rest of Internet” via AS 20

Load Sharing – Outbound Traffic (Partial routes)



Load Sharing – Outbound Traffic (Partial routes)

- R1 configuration:

```
router bgp 17821
neighbor 20.20.20.1 remote-as 20
neighbor 20.20.20.1 prefix-list DEF in
!
ip prefix-list DEF permit 0.0.0.0/0
!
```


Load Sharing – Outbound Traffic (Partial routes)

- R2 config:

```
router bgp 17821
neighbor 30.30.30.1 remote-as 30
neighbor 30.30.30.1 filter-list 30 in
neighbor 30.30.30.1 prefix-list ALL in
neighbor 30.30.30.1 route-map DEF-LOW in
!
ip prefix-list DEF permit 0.0.0.0/0
!
ip prefix-list ALL deny <bogons>
ip prefix-list ALL permit 0.0.0.0/0 le 24
!
ip as-path access-list 30 permit ^(30_)+$
ip as-path access-list 30 permit ^(30_)+_[0-9]+$
!
route-map DEF-LOW permit 10
  match ip address prefix-list DEF
  set local-preference 50
route-map DEF-LOW permit 20
```

Filter inbound routes with
AS-PATH ACL using filter-list

Accept full internet feed
except bogon routes

Purely for redundancy
(if path via AS 20 fails)

Accept routes local to and
received from AS30
(AS-path prepend included)

routes from AS30 and its
direct neighbor ASes

Low-pref default route

Using Communities

- Community attribute provides greater flexibility for traffic shaping than prefix-list
 - Simplifies BGP configuration
 - Greater policy control
- Not sent by default to BGP peers
 - Need to explicitly specify (`neighbor x.x.x.x send-community`)
- Can carry policy information
 - Example:
 - ASN:80 (set local-pref 80)
 - ASN:1 (set as-path prepend ASN)
 - ASN:888 (set ip next-hop 192.0.2.1 - Cymru bogons)

COMMUNITY recap

- Used to group prefixes (incoming/outgoing) and apply policies to the communities
 - A prefix can belong to more than one community
- Originally a 32-bit integer
 - Represented as two 16-bit integers [**ASN:number**]
 - Works well for 2-byte ASN
- With 4-byte ASNs
 - Common to see [**private-ASN:number**]
 - RFC 8092 (BGP Large Communities): 96-bit integer
 - [32-bit ASN:32-bit:32-bit]

Well-known Communities

- Many well-known communities defined
 - <http://www.iana.org/assignments/bgp-well-known-communities/bgp-well-known-communities.xhtml>
- The most commonly used:
 - **no-export**: do not advertise/export to any eBGP peers (RFC1997)
 - only extend to the neighbor AS
 - **no-advertise**: do not advertise to any BGP peers (RFC1997)
 - **no-peer**: do not advertise to any bilateral peer (RFC3765)
 - **blackhole**: null route the prefix (RFC7999)
 - Trigger blackholing
 - Signal neighboring ASes (upstream) to drop any traffic being sent towards this prefix (victim's IP addresses)

Example – IXP Route Server Communities

0 : (IXP-AS)	Do not announce any peer
0 : (PEER-AS)	Do not announce prefixes to certain peer
(IX-AS) : (PEER-AS)	Advertise to a certain peer
(IX-AS) : (IX-AS)	Advertise prefixes to all peers

Example - Equinix Community Usage

- Default Open, except with AS10, AS20, and AS30

```
router bgp 17821
  neighbor 202.79.197.126 remote-as 24115
  neighbor 202.79.197.126 route-map eqixsg-in in
  neighbor 202.79.197.126 route-map eqixsg-out out

# set default-open community on outbound and restrict
# announcements to AS10, 20 and 30

route-map eqixsg-out permit 10
  set community 24115:24115 0:10 0:20 0:30

# reject routes received from AS 10, 20 and 30
route-map eqixsg-in deny 10
  match as-path 10

ip as-path access-list 10 permit ^(10|20|30)_
```

Using communities for filters

- Set a community when you import a route from a customer or create a static (aggregate) route
 - Use that community to control export to peers & transit
 - Don't allow peers or transits to set it though
- Now when you add a prefix on a router it will automatically get exported on other routers without updating their prefix lists

Setting Communities (IOS)

```
router bgp 17821
<output omitted>
!
address-family ipv4 unicast
  network 61.45.248.0 mask 255.255.248.0 route-map SET-COMM-AGG
  network 61.45.248.0 mask 255.255.254.0 route-map SET-COMM-4G
  network 61.45.250.0 mask 255.255.254.0 route-map SET-COMM-BB
  network 61.45.252.0 mask 255.255.254.0 route-map SET-COMM-ENT
  network 61.45.254.0 mask 255.255.254.0 route-map SET-COMM-CORP
!
ip route 61.45.248.0 255.255.248.0 null0
ip route 61.45.248.0 255.255.254.0 null0 254
ip route 61.45.250.0 255.255.254.0 null0 254
ip route 61.45.252.0 255.255.254.0 null0 254
ip route 61.45.254.0 255.255.254.0 null0 254
!
```


Setting Communities (IOS)


```
route-map SET-COMM-AGG permit 10
  set community 17821:1000
!
route-map SET-COMM-4G permit 10
  set community 17821:1101
!
route-map SET-COMM-BB permit 10
  set community 17821:1102
!
route-map SET-COMM-ENT permit 10
  set community 17821:1103
!
route-map SET-COMM-CORP permit 10
  set community 17821:1104
!
```

Grouping Communities (IOS)

- We can group communities together with community-list:

```
!  
ip community-list 20 permit 17821:1000  
ip community-list 21 permit 17821:1101  
ip community-list 22 permit 17821:1102  
ip community-list 23 permit 17821:1103  
ip community-list 24 permit 17821:1104  
!
```

Two Transits and IXP



Two Transits and IXP – IXP Router

- R3 (IXP) configuration:
 - For both incoming and outgoing traffic - IXP should be the preferred path!

```
router bgp 17821
  neighbor IX-PEERS peer-group
  neighbor 80.81.192.108 remote-as 15169
  neighbor 80.81.192.108 peer-group IX-PEERS
  neighbor 80.81.192.108 description Google
  neighbor 80.81.192.283 remote-as 20940
  neighbor 80.81.192.283 peer-group IX-PEERS
  neighbor 80.81.192.283 description Akamai
!
address-family ipv4
  neighbor IX-PEERS send-community
  neighbor IX-PEERS remove-private-as
  neighbor IX-PEERS route-map IX-IN in
  neighbor IX-PEERS route-map IX-OUT out
```

Define peer-groups for all IX peers

Add neighbors to the peer group

Define common policies applied to all neighbors on the peer-group

- Send communities
- Remove private ASNs

Apply inbound and outbound routing policies

Two Transits and IXP – IXP Router

- R3 (IXP) configuration (contd..):

```
!  
ip community-list 20 permit 17821:1000  
ip community-list 21 permit 17821:1101  
ip community-list 22 permit 17821:1102  
ip community-list 23 permit 17821:1103  
ip community-list 24 permit 17821:1104  
!  
route-map IX-IN permit 10  
  set local-preference 250  
  set community 17821:<IX-AS> add  
!  
route-map IX-OUT permit 10  
  match community 20 21 22 23 24  
  set metric 10  
!
```

Define the communities

High-pref routes received from
IX peers
(outbound traffic via IX)

Define a community for all routes
learned via IXP

Send all our prefixes
(aggregates and sub-aggregates)

Set lower MED for all routes sent
to IX peers
(inbound traffic via IX)

Location info via communities

- For Transit/Upstream:
 - Tier-1 ISPs (or ISPs who are run properly) use communities to group their regional prefixes
 - Filter based on those to shape **outbound traffic** to Internet!
 - Ex: receive US routes from one ISP, and Europe routes from the other
 - Example:
 - <https://www.us.ntt.net/support/policy/routing.cfm>
 - NTT US – 2914:3000
 - NTT Europe – 2914:3200
 - NTT Asia – 2914:3400
 - NTT South America – 2914:3600

Balancing between two transits

- For Inbound traffic:
 - We can use our sub-prefixes to balance incoming traffic
 - Ex: Advertise half of our routes to one, and the other half to the other
 - keep playing until we reach symmetry!
 - But remember to announce the aggregates to both (REDUNDANCY!)
 - If very small/more specifics, use “no-export” to avoid polluting the global routing table
 - And avoid the wrath of “network police” ☺

Two Transits – Different policies

- R1 configuration:
 - Let us assume NTT (AS2914) as transit here

```
router bgp 17821
  neighbor 29.29.29.1 remote-as 2914
  neighbor 29.29.29.1 description eBGP with NTT
!
address-family ipv4
  neighbor 29.29.29.1 send-community
  neighbor 29.29.29.1 route-map NTT-IN in
  neighbor 29.29.29.1 route-map NTT-OUT out
!
! We want Asia, US and SA routes
ip community-list 1 permit 2914:3000 !US
ip community-list 1 permit 2914:3400 !AS
ip community-list 1 permit 2914:3600 !SA
ip community-list 2 permit 2914:3200 !EU
```

- Send communities
- Apply inbound and outbound routing policies

- Define communities for NTT global routes
- In this example, we will source US and Asia routes from NTT

Two Transits – Different policies

- R1 configuration (contd..):

```
route-map NTT-IN permit 10
  match community 1
  set local-preference 210
route-map NTT-IN permit 20
  set local-preference 50
!
route-map NTT-OUT permit 10
  match community 20
  match community 21
  match community 22
!
```

Route-map to influence outbound traffic

- Set higher LOCAL_PREF for US, Asia, and SA routes (outbound traffic)
- **Still lower than IX!**

Lower LOCAL_PREF for EU/rest of routes (will prefer the second ISP, but available if that link fails)

Route-map to influence inbound traffic

- Send our aggregate (in case ISP2 fails)
- And half of our sub-prefixes (*can apply no-export)

Two Transits – Different policies

- R2 configuration:
 - Let us assume Zayo (AS6461) as transit here

```
router bgp 17821
  neighbor 64.64.64.1 remote-as 6461
  neighbor 64.64.64.1 description eBGP with Zayo
!
address-family ipv4
  neighbor 64.64.64.1 send-community
  neighbor 64.64.64.1 route-map ZAYO-IN in
  neighbor 64.64.64.1 route-map ZAYO-OUT out
!
! Zayo Europe routes
ip community-list 3 permit 6461:5996
ip community-list 3 permit 6461:5998
ip community-list 3 permit 6461:5999
! Zayo Global routes
ip community-list 4 permit 6461:5997
```

- Send communities
- Apply inbound and outbound routing policies

- Define communities for Zayo global routes
- In this example, we will source EU routes from Zayo

Two Transits – Different policies

- R2 configuration (contd..):

```
route-map ZAYO-IN permit 10
  match community 3
  set local-preference 210
route-map ZAYO-IN permit 20
  set local-preference 50
!
route-map ZAYO-OUT permit 10
  match community 20
  match community 23
  match community 24
!
```

Route-map to influence outbound traffic

- Set higher LOCAL_PREF for EU routes (outbound traffic)
- **Still lower than IX!**

Lower LOCAL_PREF for global routes (NTT is preferred, but will work if that link fails)

Route-map to influence inbound traffic

- Send our aggregate (in case ISP1 fails), and
- other second-half of our sub-prefixes (*can apply no-export)

Acknowledgement:

- Philip Smith
- Cisco Systems



Questions

